

# μInference: Research platform for a micro inference server for SL5

SL5 Security Task Force

# What is the smallest inference stack we can build?

*From bootloader to inference in minimal lines of code*

# Current Inference Stack Complexity

## Traditional ML Infrastructure Stack

Application Layer	~100K-1M LOC
ML Frameworks (PyTorch, TF)	~1-2M LOC
CUDA/ROCm/GPU Drivers	~2-5M LOC
Container Runtime	~500K LOC
Kubernetes/Orchestration	~2M LOC
Linux Kernel (full)	~20-30M LOC
Bootloader (GRUB)	~300K LOC

**Total Estimate: 25-40M LOC + firmware**

# μInference Stack

## Our minimal implementation

Inference (llama2.c)
Init System (custom)
Userspace (BusyBox)
libc (musl)
Linux Kernel (minimal)
Bootloader (Limine)

4,000 LOC

**Total: ~1.3M LOC** (*3% of traditional stack*)

# What we removed from the kernel?

- Network stack
- Filesystems (except tmpfs)
- All drivers except essential for IO

# Why Limine?

- Modern UEFI/BIOS support
- Minimal configuration
- Fast boot times
- No unnecessary features

# Why musl + BusyBox?

- **musl libc (100K LOC)**
  - 10x smaller than glibc
  - Static linking friendly
  - Clean, modern C implementation
  - No legacy baggage
- **BusyBox (200K LOC)**
  - 300+ Unix utilities in one binary
  - Configurable feature set
  - 1MB static binary
  - Replaces coreutils, util-linux, etc.

# Why llama2.c?

- Complete inference in 4,000 lines
- Pure C implementation (and easily portable to Rust with LLMs!)
- Transformer architecture
- BPE tokenizer
- Temperature sampling
- Top-p sampling
- No dependencies
- Runs models efficiently



# Achieved Metrics

µInference by the Numbers

Metric	Value
Total LOC	~1.3M
ISO Size	~50MB
Boot Time	<1 seconds
RAM Usage	512MB
Inference Speed	~444 tokens/sec in consumer hardware

# Reduction Ratios

- **95%** smaller than typical Linux distro
- **97%** fewer dependencies
- **99%** less attack surface

# How Much Smaller Can We Go?

- **Phase 1:**
  - **Kernel Diet (Target: 500K LOC)**
    - Custom minimal kernel config
    - Remove more subsystems
    - Compile-time optimization
- **Phase 2:**
  - **Unikernel Approach (Target: 100K LOC)**
    - Remove userspace
    - Direct hardware access
- **Phase 3:**
  - **Bare Metal (Target: 10K LOC)**
  - No OS, just bootloader + inference
  - Custom memory management
  - Direct hardware control

# Scaling Up Capabilities

## llama.cpp Integration

- 4-bit quantization support
- GPU acceleration (Vulkan)
- ~50K additional LOC
- 10-100x performance gain
- **GPU Stack**

# Security Through Minimalism

## Attack Surface Reduction

- 97% less code = 97% fewer bugs
- No network stack = no remote exploits (There is plenty of IO options)
- No filesystem = no persistence
- Read-only system = immutable

## Security Features Possible

- Measured boot with TPM
- Memory encryption
- Secure enclave execution
- Formal verification (small enough, see seL4)

# Where Minimal Inference Matters

- Air-gapped systems
- High-security facilities
- Compliance-heavy industries
- Research sandboxes

# What This Proves

## **1. Frontier labs don't require millions of LOC**

- Core inference is surprisingly simple
- Complexity is in the ecosystem

## **2. Security through simplicity is achievable**

- Small enough to audit
- Small enough to formally verify

## **3. Edge ML is practical today**

- Runs on minimal hardware
- No cloud dependency

# Research Directions

- **Immediate:**
  - Add llama.cpp backend
  - Implement secure boot
- **Medium-term:**
  - Unikernel design
  - Deliberate Hardware stack
  - Formal verification
- **Long-term:**
  - ???



# Key Takeaways

- ✓ **1.3M LOC total** (vs 25-40M traditional)
- ✓ **50MB** complete system
- ✓ **Fully functional** LLM inference
- ✓ **Orders of magnitude** simpler
- ✓ **Auditable** and potentially **verifiable**